



CAPSTONE PROJECT BY TEAM WHAT WATTS

A DATA SCIENCE APPROACH TO FORECAST ELECTRICITY CONSUMPTION IN AUSTRALIA

(Nee) Jittinun Trairattanasirikul - (z5281789) (Ruhul) Md Ruhul Amin Sarker -
(z5275314) Peter Morian - (z5017159) James Cleaver - (z5283034)

School of Mathematics and Statistics
UNSW Sydney

June 2021

SUBMITTED IN PARTIAL FULFILMENT OF THE REQUIREMENTS OF
THE CAPSTONE COURSE ZZSC9020

Plagiarism statement

I declare that this thesis is my own work, except where acknowledged, and has not been submitted for academic credit elsewhere.

I acknowledge that the assessor of this thesis may, for the purpose of assessing it:

- Reproduce it and provide a copy to another member of the University; and/or,
- Communicate a copy of it to a plagiarism checking service (which may then retain a copy of it on its database for the purpose of future plagiarism checking).

I certify that I have read and understood the University Rules in respect of Student Academic Misconduct, and am aware of any potential plagiarism penalties which may apply.

By signing this declaration I am agreeing to the statements and conditions above.

Signed: _____ Date: _____

Signed: _____ Date: _____

Signed: _____ Date: _____

Signed: _____ Date: _____

Disclaimer

The methods, findings and recommendations in this report, as well as the source code used to produce the results in this report, are only to be used for academic purposes. Unless express prior written consent is provided by all authors involved, no part of this report and the associated source code is to be used in the commercial setting.

Abstract

Energy demand forecasting is the process of predicting the energy requirements of a jurisdiction, and is used to help the market match the electrical supply to the electrical demand to keep the electrical grid stable. The accuracy of these forecasts is important to enable the various electrical suppliers to be able to respond in a timely fashion to ensure the continual supply of stable electricity. Rooftop solar provides electricity to individual houses which removes them from current calculations of both the forecast demand, and the total demand. Current models in Australia do not include the rooftop solar into their modeling and forecasting. In this report we augment the existing forecast data with roof top solar data, temperature data, and public holiday data. We then assess four different models in an effort to create a more accurate forecasting of total demand, as well as reduce the magnitude of the larger discrepancies in the current forecast data.

We show that the XGBoost model is the highest performer with a 9.1% increase in accuracy compared to the current modeling used by the AEMO. The XGBoost model had the features identified using VIF and then used the hyperparameters of 600 estimators and linear regression as the inputs. In addition we show that in the event of an anomalous event (an event that causes a large discrepancy between total demand and forecast demand, for example large bushfires) then changing to an ARIMA model for a short period would provide a better result.

In addition to creating a more accurate model, we also describe how the model can be successfully used. Whilst our report does not include a detailed analysis of how to predict the following day rooftop solar data, we show possible methods for it to be implemented with either theoretical calculations or by using a different modeling process.

Contents

Chapter 1	Introduction	1
Chapter 2	Literature Review	3
Chapter 3	Material and Methods	5
3.1	Software	5
3.2	Description of the Data	6
3.3	Pre-processing Steps	8
3.3.1	Total electricity demand data	9
3.3.2	Importing total electricity demand data	9
3.3.3	Forecast demand data	9
3.3.4	Importing forecast demand data	9
3.3.5	Air temperature data	9
3.3.6	Importing air temperature data	9
3.3.7	NSW Public Holiday Data	9
3.3.8	Importing public holiday data	9
3.3.9	Rooftop Solar PV Data	9
3.3.10	Importing Rooftop Solar PV Data	9
3.4	Data Cleaning	10
3.4.1	Cleaning total electricity demand data	10
3.4.2	Cleaning forecast demand data	11
3.4.3	Cleaning air temperature data	11
3.4.4	Cleaning public holiday data	11
3.4.5	Cleaning solar PV data	11
3.4.6	Adding features to forecast demand data	11
3.4.7	Adding features to Temperature data	11
3.4.8	Merging All DataFrames	12
3.5	Assumptions	12
3.6	Modelling Methods	13
Chapter 4	Exploratory Data Analysis	14
Chapter 5	Analysis and Results	19
5.1	ARIMA - Baseline Model	19
5.2	Neural Network MLPRegressor	20
5.3	XGBoost Regressor	21
5.4	LSTM	22
Chapter 6	Discussion	24

Chapter 7	Conclusion and Further Issues	28
Appendix		33
	Glossary	34
	Tables	35
	Additional graphs	36

CHAPTER 1

Introduction

The AEMO is the organisation that forecasts electrical demand in NSW[1]. Since 2010 the forecast demand has shown an RMSE of 85.87, and has had 1161 number of observations with an error greater than 500. We have used RMSE as the accuracy check due to its sensitivity to outliers. The energy industry is at its most vulnerable when there is a large difference between the total demand, and the forecast demand, as it is in these events when energy suppliers need to be able to respond quickly. If this information is more accurate, the energy suppliers and the market regulator are able to be better prepared to respond to energy demand.

The purpose of this project is to create a model that can predict the demand for electricity to a higher degree of accuracy than that of the current methods used by the AEMO, and to provide more accurate forecasts of days that have a high absolute difference between forecast demand and total demand[2].

The hypothesis that we are investigating is that the gap between forecasted and total demand is partially explained by accounting for solar panels and their usage throughout the day. Roof top solar provides energy directly to the household, thus, as the number of household roof top solar installations increases, the individual demand for those houses is removed from the energy grid. We intend to use temperature data, PV data and public holiday as additional inputs into multiple models to see which gives the best forecast model. Large Solar systems such as those at greater than 30MW have been excluded as they are part of the energy market and included in the forecast and total demand equations[3].

The models we have used are MLPRegressor, LSTM, XGBOOST and ARIMA. The accuracy checks we are using to compare the models are RMSE for all the predicted values compared to the total demand, as well as the RMSE for the subset of data that has an absolute difference between total demand forecast demand of greater than 500. By taking into account solar panel usage, we intend on building a model that can predict electricity demand with more certainty. Our expectations is that the model will be used to improve the current accuracy of the data that is being provided by the AEMO.

The system would work by taking the existing forecast and augmenting the data with public holiday data, temperature of blacktown airport data, and PV data. Forecast of PV data and temperature data would be included for the following day. Figure 1.1 visualises how the system will work.

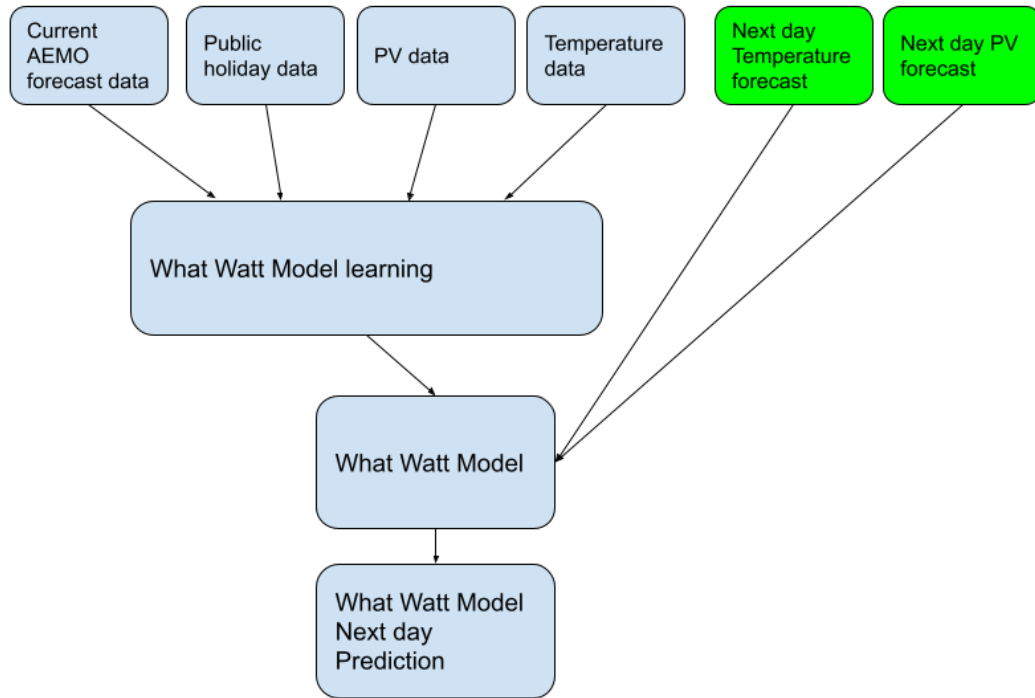


Figure 1.1: Context diagram of the proposed energy forecasting model

The historical data is provided to train the model. The forecast temperature, and PV data, and the forecast demand data would be added for the model to make a prediction.

CHAPTER 2

Literature Review

In 1879 it was realised that the construction of the Garden Palace in the domain in Sydney was unlikely to be completed in time, as such the Government at the time decided to work at night. Arc lighting was chosen to illuminate the site [4]. The demand forecasting process was slow, as was the electricity supply it required several electrical generators to be procured and shipped from England[5].

As noted by Brady significant events in Australian history have been as a result of electricity demand, these include metallurgical industry in 1914, The Snowy Hydro scheme beginning in 1949 [4] these have had significant impact in shaping Australia.

Bently describes how the oil crisis in 1971 led to the need for better and more accurate energy forecasting [6]. It was also around this time that it was recognised that the speed of forecasting was important to manage not just the long term demand, but also the change in demand across the duration of a day.

In 1997 NSW opted to privatise the energy industry[7] and systems needed to be put into place to ensure that the market was able to match supply and demand of electricity. Electricity requires the demand to equal supply for the electrical grid to remain operational, if there is an increase or decrease in the supply compared to the demand, then the frequency of the electricity supply moves away from the Australian standard of 50hz, and if moving too far away may result in a catastrophic failure of the electrical grid.

Up until 2007 the majority of electrical supply was provided by the main market players, with coal and the commercial green energy suppliers providing to the market. In 2007 there was a move towards rooftop solar, which is a system whereby individual households were able to generate electricity that would be supplied into the grid. The systems in place to match the electrical supply with the electrical demand, did not include the rooftop solar. This is due to the fact that the rooftop solar provides electricity to the house it is on, and that house's energy demand is not included into the forecast that is available to be sold to the electrical grid.

The rooftop solar has continued to grow, in 2007 it was 0.2% of the household roofs, whereas in 2020 it had moved to 20% of houses having rooftop solar, and using photovoltaic cells to generate local electricity[3].

As the uptake and increase of rooftop solar has increased, this has resulted in greater difficulty in modeling forecast demand[8]. The fact that an increasing number of houses are providing their own electricity for periods of the day means that the variance in electricity forecast reliability decreases. Currently the electricity used by domestic households for heating and cooling is 32%-33% of total domestic energy use and is expected to remain steady at around that number[9]. This heating

and cooling is driven by the air temperature, and is the reason why temperature is a key focus area in predicting energy demand.

The modeling methods of energy demand has changed in line with the changes in energy usage. When the Snowy Hydro first came in the majority of the electricity demand was from industry, and highly predictable, there was less need for highly sophisticated modeling, however as the domestic appliances, heating and cooking all moved to be driven by electricity the need to model in a more granular way became essential. In the 1970s the modeling was based on simple time linear equations[10][11]. Since then modeling has improved to move towards more robust linear models, and then through to time series modeling, and then more recently neural networks[12]. Modeling has included linear neural networks and LSTM neural networks and MLPRegressor networks.

The focus of the modeling has been 2 fold, one to increase the accuracy of the forecast, but just as importantly to decrease the maximum errors, when the forecast model provides a result that turns out to be vastly different to the actual demand, then there is significant strain on the electrical grid. When there is a large difference between the forecast and the actual demand it is necessary for the electrical suppliers to be able to respond rapidly to ensure the stability of the frequency in the network. In terms of speed to respond, Coal takes the longest, pumped hydro and hydro the next, solar and wind (assuming conditions are favourable), faster still and the quickest is the new grid connected battery farms that can respond in microseconds.

CHAPTER 3

Material and Methods

3.1 Software

All models in this report (except for Time Series) have been built using Python via Google Colab, Table 3.1 shows the libraries and machine learning models used.

Library	Description
pandas	Open source data analysis and manipulation tool.
numpy	Fundamental package for scientific computing with Python.
seaborn	Statistical data visualization.
matplotlib	Creating static, animated, and interactive visualizations.
plotnine	Grammar of graphic, based on ggplot2
sklearn	Simple and efficient tools for predictive data analysis.
plotly	interactive, publication-quality graphs.
XGBoost	Efficient and portable optimized distributed gradient boosting.
statsmodels	Estimation of many different statistical models.
scipy	User-friendly and efficient numerical routine.
keras	Deep learning framework.
MLPRegressor	Multi-layer Perceptron regressor, optimises the squared-loss using LBFGS or stochastic gradient descent.
LSTM	Long Short-Term Memory, a model capable to learn based on order dependance in sequence.

Table 3.1: python libraries and models used in the project

The Time Series model was built using R & RStudio[13], with the libraries in Table 3.2. All scripts used to produce the models and results in this report are available on GitHub.

Library	Description
data.table	A data format used in R.
dplyr	Used for data manipulation and transformations.
forecast	Methods Tools to conduct time series analysis.
lubridate	Tools for transforming data/time variables.
psych	Produces summary statistics of data.

Table 3.2: R libraries used in the project

For EDA, Power BI was used to produce visuals that explain trends and relationships between variables. Data visualisations produced by Power BI are included throughout this report, and the PowerBI code used is available in GitHub.

This report was produced using Rmarkdown and MS SQL Server was used to process the PV data to use for the model[14].

3.2 Description of the Data

The data provided consisted of 3 csv files, all the files are rectangular. The file forecastdemand_nsw.csv is 739.63 MB in size and consists of 6 columns and 10,906,019 rows. The type of data is shown in Table 3.3. This data has been generated by the AEMO[15].

Column Name	Python Data Type	Variable type
PREDISPATCHSEQNO	Int64	Discrete
REGIONID	object	Categorical 1 Value
PERIODID	int64	Categorical
FORECASTDEMAND	float64	Continuous
LASTCHANGED	object	Discrete
DATETIME	object	Discrete

Table 3.3: Forecast demand data properties

The file totaldemand_nsw.csv is 5.80 MB in size and consists of 3 columns. The type of data is shown in Table 3.4

Column Name	Python Data Type	Variable type
Column Name	Python Data Type	Variable type
DATETIME	object	Discrete
TOTALDEMAND	float64	Continuous
REGIONID	object	Categorical 1 Value

Table 3.4: Total demand data properties

The file temperature_nsw.csv is 6.877450MB in size and consists of 3 columns. The type of data is shown in Table 3.5

Column Name	Python Data Type	Variable type
LOCATION	object	Categorical 1 Value
DATETIME	object	Continuous
TEMPERATURE	object	Continuous

Table 3.5: Temperature data properties

The additional data we sourced was PV data, and public holiday data. Both files are rectangular, the file “NSW Public Holidays.csv” is 8713B in size and consists of 5 columns and 135 entries, this was sourced from the Digital Transformation Agency [16] and then had additional lines added taken from online calendars. The type of data is shown in Table 3.6

Column Name	Python Data Type	Variable type
Year_Date	Datetime	Discrete
Day	object	Categorical
Holiday	object	Categorical
Type	object	Categorical
Other	object	Categorical

Table 3.6: Public Holiday data properties

We downloaded PV performance data for different NSW regions and PV installation size for NSW and processed those files by using MS SQL[17]. The type of data for the sites is shown in Table 3.7 and for installations in Table 3.8

Column Name	MS SQL Data Type	Variable type
Location	Varchar(200)	Categorical 1 value
Timestamp	Datetime	Discrete
Irradiance(W/m2)	Decimal(28,3)	Continuous
System temperature (C)	Decimal(28,3)	Continuous
Nearest BOM station temperature (C)	Decimal(28,3)	Continuous
PV Yield (kWh)	Decimal(28,3)	Continuous

Table 3.7: NSW Regions Photo Voltaic data properties

Column Name	MS SQL Data Type	Variable type
Month	Varchar(50)	Categorical 1 value
lt2.5kW	int	Continuous
2.5–4.5	int	Continuous
4.5–6.5	int	Continuous
6.5–9.5	int	Continuous
9.5–14	int	Continuous
14–25	int	Continuous
25–50	int	Continuous
50–100	int	Continuous
100kW–5MW	int	Continuous
5MW–30MW	int	Continuous
30+ MW	int	Continuous
Total Size (KW)	int	Continuous

Table 3.8: NSW Regions Photo Voltaic data properties

The file PVGenerationNSWACT-30+MW-Excluded-2012_032021.csv is 3.37MB in size and consists of 162144 rows and consists of 2 columns. The type of data is shown in Table 3.9 The file was the final dataset extracted from MS SQL.

Column Name	Python Data Type	Variable type
DateTime	object	Discrete
PVGeneration(MW)	int64	Continuous

Table 3.9: Photo Voltaic generation data properties

3.3 Pre-processing Steps

The following three data sets containing historical data from 01 January 2010 were shared through a GitHub repository. These files were the ones provided as initial input to the electricity modeling problem.

- Total electricity demand
- Forecast demand
- Air Temperature

We also used the following two external data sets for this project.

- NSW public holiday data
- Rooftop Solar PV data

To work in Python the GitHub repository was cloned and a data folder was created. There were two files for the forecast demand which merged into one file. All three data files were unzipped and stored as .csv in the data folder.

3.3.1 Total electricity demand data

The Total electricity demand in megawatt is half-hourly increments demand for New South Wales. This data is sourced from the Market Management System database, which is published by the market operator from the National Electricity Market (NEM) system.

3.3.2 Importing total electricity demand data

The data from the `totaldemand_nsw.csv` was stored into a dataframe.

3.3.3 Forecast demand data

The Forecast demand data is the half-hourly increments of electricity demand in megawatt for New South Wales. This data is also sourced from the Market Management System database.

3.3.4 Importing forecast demand data

The data from the `forecastdemand_nsw.csv` was stored into a dataframe.

3.3.5 Air temperature data

The air temperature data for New South Wales was measured from the Bankstown Airport weather station. This data is sourced from the Australian Data Archive for Meteorology.

3.3.6 Importing air temperature data

The data from the `temperature_nsw.csv` was stored into a dataframe `df_temp`.

3.3.7 NSW Public Holiday Data

The NSW public holiday data was taken from Digital Transformation Agency for 2014 onwards, and then manually updated with the 2010 through to 2014 based on searching old calendars[16].

3.3.8 Importing public holiday data

The public holiday data was loaded into a DataFrame `df_publicholidays` from `NSW Public Holidays.csv`.

3.3.9 Rooftop Solar PV Data

Following two different data sources related to solar PV were downloaded from (Australian PV Institute (APVI) Solar Map) for January 2010 to March 2021 period [17].

- Solar PV performance data for several regions of NSW
 - Canterbury (NSW)
 - Sutherland (NSW)
 - St Ives (NSW)
 - Newtown (NSW)
- Cumulative PV installation size data for NSW and ACT regions

3.3.10 Importing Rooftop Solar PV Data

We used MS SQL Server database to store and process the data to prepare the final PV Data set to load into Python DataFrame for the model.

- Both solar PV data sets were loaded in MS SQL Server database by using SQL Server Import and Export Wizard.

- As the size of a single PV unit for different locations are different (see Table 3.10), we normalised the PV Yield (kWh) value for a PV unit size of 5kWp for all locations.

Region	Single PV system size (kWp)
Canterbury (NSW)	3.06
Sutherland (NSW)	5.04
St Ives (NSW)	5.1
Newtown (NSW)	5.1

Table 3.10: Photo Voltaic generation region and single system size

- Prepared a dataset by using only Canterbury PV performance data as this region has been considered as a representative of NSW solar PV data (see Section 3.5 for detailed explanation).
- The missing dates of Canterbury were filled by using either Sutherland or St Ives or Newtown data.
- Canterbury solar PV performance data is for every hour interval. We converted the data to make a 30 minutes interval of solar power production by halving the hourly data.
- PV installation of size 30+MW excluded from the dataset during calculating the total cumulative size of PV installation for NSW for every month[18] we normalised the PV Yield (kWh) value for a PV unit size of 5kWp for all locations (see Section 3.5 for details).
- Final dataset was exported as a csv file.
- The final rooftop solar PV generation data was loaded into a DataFrame.
- We prepared a dataset for every 30 minutes interval by using solar power generation for NSW by multiplying Canterbury data with the total capacity of NSW PV for that date and time. The PV capacity data was based on monthly totals and we assumed a linear rollout of the PV across that month (this had small impact on the total numbers but was a reliable method).

3.4 Data Cleaning

Data cleaning is an essential part of every machine learning project to get better and accurate models. The following summarizes the data cleaning steps that we undertook.

3.4.1 Cleaning total electricity demand data

- The total demand data was time series data with values of total demand for every 30 minutes, this was treated as the ground truth and the sampling period that the rest of the data would be aligned to.
- The DATETIME column was imported as an Object which was converted to python Datetime.
- we set the DataFrame index by using DATETIME

3.4.2 *Cleaning forecast demand data*

- The DATETIME column of forecast DataFrame was imported as an Object which was converted to python Datetime.
- we set the DataFrame index by using DATETIME
- The Forecast Demand data had one or more values for each of the 30 minutes time periods. As the later values represented the most recent output from the NEMS forecast demand, the most recent forecast demand was considered ,as this was the output from the model which had the most up-to-date input data.

3.4.3 *Cleaning air temperature data*

- The DATETIME column of temperature DataFrame was imported as an Object which was converted to python Datetime.
- We set the DataFrame index by using the DATETIME column.
- The temperature data was a time series that was more asynchronous and could contain multiple values for the 30 minute target time. This data was always relatively close to each other so an arithmetic mean was used to create the single value for the 30 minute period.

3.4.4 *Cleaning public holiday data*

- Renamed the column Year Date to DATETIME
- Excluded pre 2012 data from the DataFrame
- Added a flag column PUBLICHOLIDAYS
- Set the DataFrame index by using the DATETIME column.
- Dropped unused columns Day, Holiday, Type and Other
- Resampled the public holiday data for everyday
- Updated flag column PUBLICHOLIDAYS with 0 for non public holidays
- Resampled the public holidays DataFrame for 30 minutes interval

3.4.5 *Cleaning solar PV data*

- The DateTime column of solar pv DataFrame was imported as an Object which was converted to python Datetime.
- Set the DataFrame index by using DateTime
- Renamed the column DateTime to DATETIME

3.4.6 *Adding features to forecast demand data*

- Added hour, Minute, HourMinute, Month, Year, Day, DayName as extra columns
- Added season based on Month from December to February as ‘Summer’, March to May as ‘Autumn’, June to August as ‘Winter’ and September to November as ‘Spring’

3.4.7 *Adding features to Temperature data*

Added temperature classification which is grouped into 5 different ranges.

- Less than or equal 10 degree C: very low
- Between 10 and 20: low
- Between 20 and 30: high
- Between 30 and 35: very high

- More than or equal to 35: extremely high

3.4.8 Merging All DataFrames

Created a dataframe `df_final` by merging all five dataframes (`df_forecast_merge`, `df_temp_merge`, `df_totaldemand`, `df_publicholidays_merge`, `df_pv`)

One hot coding on columns which have categorical data, as we have introduced categorical data into our datasets, we were required to apply one-hot encoding techniques in order to use these features in our models. One-hot encoding expands each possible value of a categorical variable into multiple binary columns, so that these features can be interpreted by the models[19]. The following categorical variables were transformed using one-hot encoding:

- Seasonal categorical data
- 0 or 30 minute categorical data
- Temperature categorical data

3.5 Assumptions

- The forecast demand data had one or more values for each of the 30 minutes time periods. As the later values represented the most recent output from the NEMS forecast demand, the most recent forecast demand was considered as this was the output from the model which had the most up-to-date input data[20].
- Assume the PV data at Canterbury was representative of the PV data across the highest demand area of the state (the Sydney / Newcastle / Wollongong areas). The figure 3.1 shows the comparison of PV performance for Canterbury and St Ives. Figure 7.2 in the Appendix shows the comparision with Sutherland[20].



Figure 3.1: PV Performance comparison between Canterbury and St Ives generated by single 5Kwp system

- We have assumed that the 30+MW size solar PV systems are part of the Energy Grid, and participate in the energy spot market. This means that their capacity is included in the total demand / forecast, and as such should be excluded from our PV data, which is to represent supply that is not part of the spot market[18].

- The temperature at Bankstown airport is indicative of the entire state. This assumption is based on the similarity of climate zones[20] where the majority of the energy demand is from.
- To minimise the possibility of multicollinearity, variables were removed from the training dataset based on their Variance Inflation Factor (VIF). Should a variable have a high VIF, then that variable would be removed as it significantly increases the variance of parameter estimates due to high correlation with other variables. The generally accepted rule in the industry is that a variable with a VIF greater than 10 should be removed, Table 7 in the Appendix shows how the rule was applied to our variable selection criteria[21].
- We assumed that 500 was a good value for the outliers, due to the natural split in the data when viewed, this can be inferred from Figure 6.1.

3.6 Modelling Methods

Due to research such as Almalaq et al [22] and Marcjasz et al [12] We decided to use several different models to improve the RMSE comparatively against the AEMO data. The first, simple model that we chose was an ARIMA time series model, this was to help us get an understanding of the data, and what we might be able to expect from future modeling. We then chose to use a MLP Regressor model, and XGBoost model, and an LSTM model. The MLPRegressor would act as a sliding window across the data, building up a good representation of the data, the XGBoost uses a regulated Gradient Descent approach that should assist with identifying outliers, and the LSTM was chosen as it should provide the capability of understanding the previous steps, similar to ARIMA as well as pick up on the features more accurately similar to the MLPRegressor and XGBoost models.

CHAPTER 4

Exploratory Data Analysis

The first step in our exploration of the data was to visualise the target output, in this case total demand. The figure 4.1 shows the scatter plot of total demand and the figure 4.2 shows the forecast demand and the figure 4.3 shows the scatter plot of the difference between total demand and forecast demand averaged across a day. As can be seen, the data has a slow trend downwards over the year with weekly cycles. The forecast demand had a high degree of uncertainty. The years 2010-2012 had very high variance, and as a result we dropped these from the modeling early.



Figure 4.1: Total demand by datetime



Figure 4.2: Total forecast by datetime




Figure 4.3: Absolute difference between total demand and total forecast averaged across a day for 2010 to 2021

In an effort to understand how the demand changed over the years we graphed the probability density against the energy consumption, this would enable us to see the similarity in the graphs as well as to see how the peak value changed Figure

4.4 shows this information. It was observed that the unimodal distribution for each year resulted in a lower energy consumption each year. This shift is most likely be caused by the increasing installation of rooftop solar panels and specially in the 2020 the cause is the Covid-19 shutdown. With rooftop solar not included in the total demand numbers this could be explained by the addition of rooftop solar.

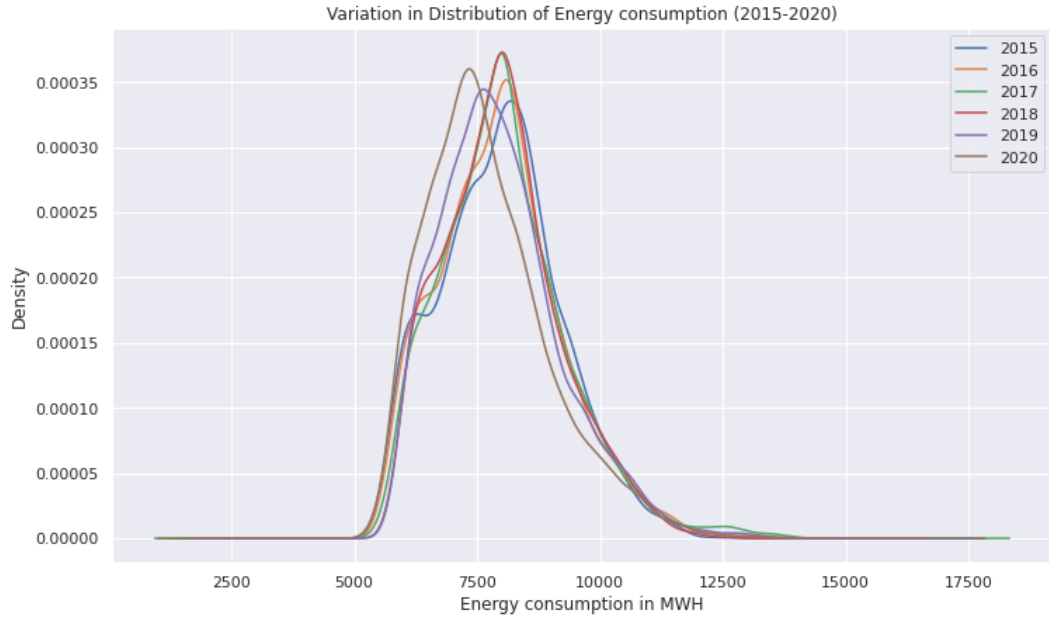


Figure 4.4: Difference between total demand and forecast vs Temperature for every year from 2015 to 2020

As we were investigation the impact of PV energy in the model, it was important to understand how the impact would be seen across the seasons. During summer PV data should show a greater impact than in the winter. To investigate this, we graphed the distribution for summer versus winter as shown in figure 4.5, the summer data was bimodal, compared to the winter which was unimodal. Whilst the spread of the data was similar between winter and summer the variation year on year increased, with 2017-2019 are bimodal in summer, which is an indication of the rooftop solar rollout. The 2020 was moving more to the unimodal type again, this could either be due to the large percentage of rooftop solar impacting the data, or due to covid social impact changing how people responded.




Figure 4.5: Variation for summer and winter of total demand for the years 2015 to 2020

In household and large corporations Ürge-Vorsatz et al show that the use of electricity is often linked to either heating or cooling [9]. To see if this was an impact to our model we did a scatter graph of total demand versus temperature. Figure 4.6 shows that the total demand is lowest around 18-19 degrees and increases the further you move away from that temperature. The data shows that this is consistent across all the years from 2015 - 2020.




Figure 4.6: Difference between total demand and Forecast vs Temperature for every year from 2015 to 2020

When looking at the overall scatter shown in Figure 4.1 it could be seen that the data followed weekly cycles. To see if this was a result in similarities between

daily information or other influences, we created a box plot of demand versus week day. Figure 4.7 shows that the week days are very similar and the weeknds separate again. Mondays did have some similarities with weekend's this is likely due to the fact that public holidays are often on a monday, and the demand on those days is more similar to the weekend. As a result of this data we chose to use weekday and public holiday data as features to include in the modeling (this was then verified by the VIF we performed).




Figure 4.7: Electricity Total Demand distribution on week days for 2015 to 2020

Our focus was in identifying how we could improve the accuracy of outlier days, as well as the accuracy of the overall model. For this we needed to have a good understanding of where the outliers are. Figure 4.8 is a box plot that shows the different quartiles for each year form 2015 through to 2020.

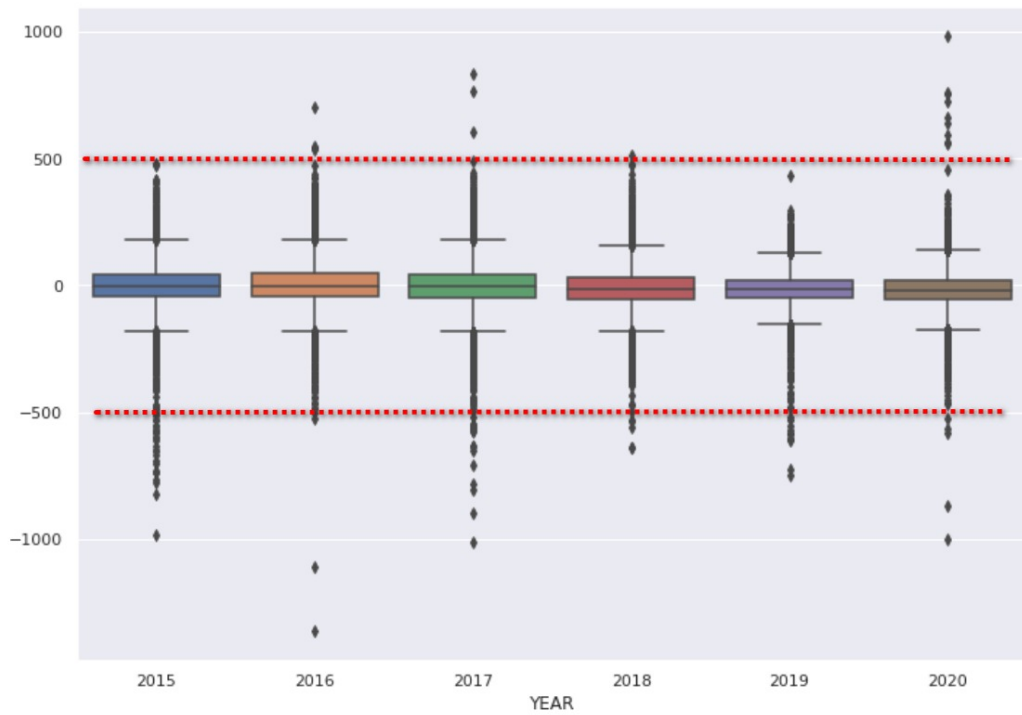


Figure 4.8: Distribution of the difference between Total demand and Forecast for every year from 2015 to 2020

The output from the rooftop solars has been increasing over the last 10 years, whilst there has been a slow decline in the total demand, Figure 4.9 indicates that the increase in PV output is increasing it does not account for the total decrease in total demand. Figure 7.1 in the Appendix visualises the growth of rooftop solar




Figure 4.9: PV Generation vs Total Demand for 2012 to 2020

CHAPTER 5

Analysis and Results

5.1 ARIMA - Baseline Model

Time series models are often used to predict future events when given data that is provided in regular & discrete time intervals[23]. It is an especially useful modelling tool when the data has a clear trend and/or seasonality components, as these predictable features can be isolated via the Box-Jenkins method, and an AutoRegressive Integrated Moving Average model (ARIMA) can be applied. Time series models are commonly used in stock price predictions, logistics management, sales forecasts and IoT devices.

As the data for total demand was provided in regular 30-minute intervals, and since there is a clear seasonality trend in the data, it was decided that a time series model would be appropriate to use when predicting future total demand.

The ARIMA model was trained on the first 80% of the data (which contains data from 1/1/12 to 16/5/19), and then tested the remaining 20% of the data. Using the “auto.arima” function in R to find the combination of p, d & q that minimises the AIC. An ARIMA(4,1,2) model was suggested to have the optimal set of parameters for the ARIMA, as shown in Figure 5.1 below [24]. Mathematically, this equation is written below as, where \hat{X}_t represents total demand at time t, and ϵ_t represents the error term at time t:

$$\hat{X}_t = 0.8243X_{t-1} + 0.8931X_{t-2} - 0.7242X_{t-3} - 0.0681X_{t-4} + \epsilon_t - 0.0446\epsilon_{t-1} - 0.9328\epsilon_{t-2}$$

The performance of the ARIMA(4,1,2) model has RMSE of 82.48 for 2020 which is higher than RMSE of the AEMO’s model which is 74.15 which indicated that the ARIMA(4,1,2) model performed poorer than the AEMO’s model for overall predictions (see Figure 5.1).




Figure 5.1: Absolute Demand Difference of AEMO’s model versus ARIMA Model for 2020

For outliers, the ARIMA(4,1,2) model has RMSE of 244.41 which is much lower than RMSE of the AEMO's model which indicated that the ARIMA(4,1,2) model performed much better for the outliers (see Figure 5.2)




Figure 5.2: Absolute Demand Difference of AEMO's model versus ARIMA Model for outliers in 2020

5.2 Neural Network MLPRegressor

The MLPRegressor model acts as a sliding window across the data, researchers such as Ozden et al. have used MLPRegressor models to forecast electricity demand, in their case it was to assist with determining locations of where renewable power supplies should be placed[25].

We used 200 nodes in hidden layers as it offers the best output after we tested using 50, 100, 200, 300, 400, 500 and 1000 nodes. We used the relu activation function because it gives output between 0 and the maximum value also Ozden et al. used it in their research. We use ADAM solver for weight optimisation as it offers the better output than SGD solver. We use alpha L2 penalty (regularisation), termination parameter as 0.001, the maximum number of iterations (max_iter) as 5000, and the tolerance for the optimisation as 0 as they offer better output than others tested.

The overall performance of the MLPRegressor model was better than some models in this paper (ARIMA, LSTM); however, XGBoost performed better than MLPRegressor. It offers the optimal output with RMSE of 70.99 comparing to RMSE of 74.15 from the AEMO's model (see Figure 5.3). The reason is that because the MLPRegressor model optimises the squared-loss using LBFGS or stochastic gradient descent.




Figure 5.3: Absolute Demand Difference of AEMO's model versus MLPRegressor for 2020

For the outliers, the MLPRegressor model outputs RMSE of 710.96 which is lower compared to RMSE of 714.42 from the AEMO's model (see Figure 5.4). For the outliers days, MLPRegressor performs better than XGBoost Regressor; however, not as good as ARIMA model which has information about unusual events (for example, due to bushfires, COVID lockdown) to be able to predict outliers correctly Figure 7.3.

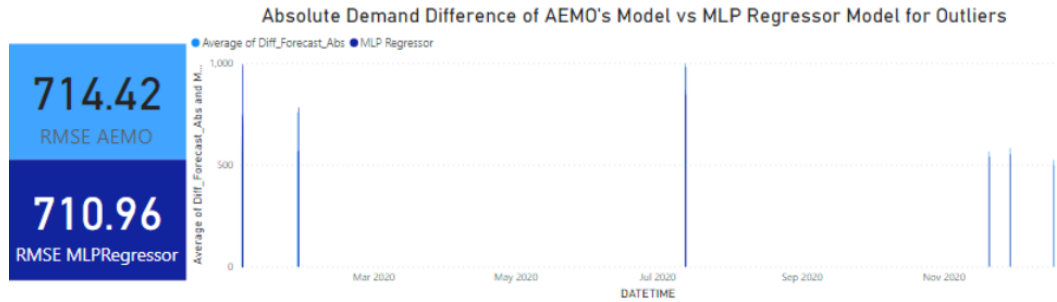


Figure 5.4: Absolute Demand Difference of AEMO's model versus MLP regressor Model for outliers in 2020

5.3 XGBoost Regressor

The XGBoost Regressor model was trained on 2 recent years of data between 2018 and 2019, including year 2021 data upto 17th March 2021 and tested on one year of data in 2020. The model calls XGBRegressor library in Python with 600 estimators, linear regression to train and test. Using 600 estimators in this model offer the best output which is the lowest RMSE as 67.35 comparing to using other number of estimators (see Figure 5.5).




Figure 5.5: Test RMSE versus XGBoost Estimator

The overall performance of the XGBoost Regressor model was greater than other models in this paper (ARIMA, and LSTM). It offers the optimal output with RMSE of 67.35 comparing to RMSE of 74.15 from the AEMO's model (see Figure 5.6). The reason is that XGBoost regressor offers optimised implementation of gradient boosting which calculates residual errors from the latest predictor and tried to fit the new predictor to these residual errors. XGBoost includes a unique split-finding algorithm to optimise trees, along with built-in regularisation that reduces overfitting. It is scalable and very fast. This model is also easy to tune with minimum parameters to tune and fast to re-train.




Figure 5.6: Absolute Demand Difference of AEMO's model versus XGBoost Regressor for 2020

For the outliers, the XGBoost Regressor model outputs RMSE of 741.26 which is higher compared to RMSE of 714.42 from the AEMO's model (see Figure 5.7). For the outliers day XGBoost Regressor performs not as well as ARIMA model as XGBoost Regressor does not have information about unusual events (for example, due to bushfires, COVID lockdown) to be able to predict outliers correctly Figure 7.3 in the appendix shows this information.




Figure 5.7: Absolute Demand Difference of AEMO's model versus XGB Regressor for Outliers in 2020

XGboost has been used by researchers such as Almalaq et al to model demand forecasts[22].

5.4 LSTM

The long short term memory models are often used to pick up non-linear behaviour in time series data sets. The LSTM has the ability to build a model that forgets information not relevant, and focuses on the long term trends as well as the short

term fluctuations. Researchers such as Mesa Jiménez et al. have used LSTM models to model demand, however they also used exponentially moving averages to help capture the peaks in the data. Due to the limitation in time, we used a naive approach to the LSTM modeling[26].

We created a LSTM that consisted of 200 neurons (tested against 50, 100, 150 and 200), we also had a network that consisted of 4 layers, 3 LSTM layers, the final with a 50% dropout to avoid overfitting, and the final being a fully connected dense layer to generate the predicted total demand. The lag of the series consisted of only a single previous value, and this naivete led to high RMSE value of 195.585, If additional time, we would have built an LSTM that had a lag of 48 to be able to have a full days history in the model. The additional layers should then have been able to pick up more of the characteristics of the data. The model was created by using 2 years worth of data, and with 2020 data as the test data. As the LSTM requires normalised data, all the data (including the one hot encoded data) was normalised to be between 0 and 1. The predicted total demand hours where then rescaled back so the data could easily be compared. The naivety of the model and the lack of time to improve its accuracy meant that we reverted to the MLPRegressor, XGBoost models, and ARIMA models.

CHAPTER 6

Discussion

To analyse the performance of the models we graphed the RMSE of each of the models to make a comparison (Figure 6.1). We found that the most difference between total demand and forecast demand from the AEMO's model is either greater or lower than 500 MWh which we consider as outliers for our analysis.




Figure 6.1: RMSE comparison of the prediction for different models

The performance of the ARIMA(4,1,2) model was not as great as the XGBoost, as it has a larger RMSE. One possible explanation for this is due to the fact that time series only takes into account the past values of a single attribute when modelling the future values of that same value. As there are a multitude of factors that can influence electricity demand, having a model that only considers one feature is not ideal. In addition to the fact that the performance of the XGBoost is superior to that of the time series, it was collectively agreed upon that we would not further investigate other ARIMA, ARIMAX or SARIMA models, as it does not meet our initial objective of building a model that accounts for factors that are currently not included in the AEMO's model. However, ARIMA is good at detecting outliers which should be used for prediction of unusual events.

The performance of the MLPRegressor model was better than ARIMA as it acts as a sliding window across the data to input 200 data points into 200 hidden

nodes at a time resulting in not having historical information to be able to detect outliers well compared to ARIMA. However, MLPRegressor didn't perform as well as XGBoost Regressor as the overall performance of the XGBoost Regressor model was greater than other models (the AEMO's model ARIMA, LSTM and MLP Regressor) as it offered the lowest RMSE. The reason is that XGBoost regressor offers optimised implementation of gradient boosting which calculates residual errors from the latest predictor and tries to fit the new predictor to these residual errors. XGBoost includes a unique split-finding algorithm to optimise trees, along with built-in regularisation that reduces overfitting. It is scalable and very fast. This model is also easy to tune with minimum parameters to tune and fast to re-train. For the outliers, XGBoost Regressor performs worse than the ARIMA model as XGBoost Regressor does not have information about history and trends which caused by some unusual events (for example, due to bushfires, COVID lockdown in 2020) to be able to predict outliers correctly.

Due to the limited outliers in the train dataset, we further investigate by training data with the more outliers from 2012 onwards. We discovered that the more outliers to train, the better the model as outliers are a sparse dataset contributing to accuracy of predictions. Due to its complexity and its ability to retain prior information in predicting future values, our initial expectations were that the LSTM model would have produced the lowest RMSE. Thus it was rather surprising that the LSTM model had one of the worst performing outcomes out of all the models tested. Whilst further work can be done to improve the performance of the LSTM (further hyperparameter optimisation, more epochs, and a greater lag etc.), it is unlikely that there would be significant gain in predictability in the timeframe we had.

To visualise the performance of the models, we compared the RMSE for the model, as well as for the outliers. Figure 6.2 Shows that comparison, it clearly shows the benefit of the ARIMA for outlier events, as well as the benefit of the XGBoost for normal day to day operations.




Figure 6.2: RMSE comparison of the prediction for all data and the outlier subset for the selected models

In addition, to visualise the accuracy of each of the models we used a binning technique to graph the difference between total demand and predicted demand in bin sizes of 100, it was interesting to note, that very few of the models wever got the predictions correct, and in fact the most common values of diffeence where in the range of 100-200. Figure 6.3 shows this.




Figure 6.3: Comparison of binned values of total demand - predicted demand for each model

Whilst our models where able to increase the accuracy of the both the outliers and the overall model, it is important that the capability we have created is easily able to be used. This capability is dependent on several components, it requires the input of the temperature data, the current forecast data, the public holiday

data, as well as the PV data. In the current modeling that the AEMO does, they have the forecast, public holiday, and the temperature data. This means the only new data that they need to build is the PV data. The AEMO only used the public holiday data to differentiate between working day, and non working day, we have used more granular categories to build a more accurate model.

The PV data was created by using a single 5.1kW system as the baseline, this was then normalised to a per MW value. The total MW usage across the state was then used to scale this up to current usage. The total MW values are monthly based forecasts, and this data could be refreshed each month, and used for future forecasts. The daily data is more challenging to predict, whilst not in the scope of the project, we have identified several ways this could be done. The theoretical amount of PV data can be calculated from the latitude, longitude and the air mass index at the time of the forecast. This data follows a similar curve to that which has been seen by the single system. This theoretical maximum would then be multiplied by a normalisation factor calculated by comparing several days worth of forecasts. The PV data could be predicted based on the BOMs UV data or similar, and a model built up and compared to the current PV data. Alternative research by Maleki et al. provides a review of several other techniques if more precision is required[27].

The limits of the capability of our model is also the ability to predict the outlier days. The ARIMA model was the best at calculating the outliers, this is because for most of the extreme events where the difference between forecast demand and the observed total demand did not correlate well with either temperature or PV data. Whilst they nearly all occurred on either extremely hot, or very cold days, there was also many extremely hot or very cold days that the model forecasted accurately. We believe the reason for this is that the outlier events were caused by other social phenomena, or impactful events. For instance the outliers on the 4th January 2020 were likely caused by the bushfires and the failure of one of the electrical power stations, similarly the outlier event on the 13th July 2020 was likely caused by COVID lockdown in Sydney. Our recommendation would be that in the event of a large discrepancy between the XGBoost forecast and the observed total demand, then changing to the arima model for a period of time would be preferable.

CHAPTER 7

Conclusion and Further Issues

The modeling that we have done shows that our system has the capability of increasing the accuracy of AEMO's modeling. We have shown that the use of PV data and by more granular categorisation of public holiday and weekday data it is possible to build a more accurate model. The XGBoost model that we generated was able to provide a 9.1% increase in performance when compared to the current AEMO's prediction.

When comparing between models, we had expected that the LSTM model would be the most accurate, however due to the naivete of our model this did not result in an accurate prediction of total demand. The MLPRegressor model was only slightly worse than the final XGBoost model, this is to be expected as they both rely on similar underpinning optimisations. The ARIMA model was surprisingly accurate for the outlier events, this is likely due to the fact that the lack of correlation between the data sources used and outliers means that the previous values is a better indicator of total value than other data observations.

For our model to be successfully implemented it requires AEMO to start collecting PV data, and PV forecast data. It requires both the PV data for a single system (or theoretical calculation) multiplied by the current installed base. It is likely that as PV rollout continues across the state, this will become an increasing factor in both detecting and managing anomalous events in the future.

The best implementation of our model would be done with a parallel implementation of both the ARIMA and the XGBoost model, whilst the absolute difference between forecast demand and total demand was less than 500 the XGBoost model should be used, however, when the absolute difference moved outside of this value, then it should use the ARIMA models output. Due to time constraints, we were unable to model the duration of time to run the ARIMA model, however, it is likely that this should only be for a period of a couple of hours. It should be noted that attempting to use a less than absolute difference of 500 when end up with an unstable oscillation between the models, hence further research is required to accurately build this implementation.

It was also noted that whilst the PV Data does increase the accuracy of the model, the outlier events do not strongly correlate with either PV, temperature or public holiday data, it is likely that an additional data set would be required to predict outlier events. During the process of the project we identified that both COVID and bushfires were likely to be the cause of outlier events.

As shown in the literature review as technologies have changed, so has the need for the electrical industry to respond. The need for accurate and timely forecasts of demand is becoming increasingly challenging. Over the coming years energy storage

devices such as large capacity batteries and pumped hydro as energy storage will change the landscape of energy demand forecasting. Future research should include identifying and creating additional data sources in this area.

References

- [1] Data model reports.
URL https://visualisations.aemo.com.au/aemo/di-help/Content/Data_Model/MMS_Data_Model.htm?TocPath=____8
- [2] S. C. Johnson, D. J. Papageorgiou, D. S. Mallapragada, T. A. Deetjen, J. D. Rhodes, M. E. Webber, [Evaluating rotational inertia as a component of grid reliability with high penetrations of variable renewable energy](#), *Energy* 180 (2019) 258–271. doi:<https://doi.org/10.1016/j.energy.2019.04.216>.
URL <https://www.sciencedirect.com/science/article/pii/S0360544219308564>
- [3] AEMO, [Projections for distributed energy resources -solar pv and stationary energy battery systems report for aemo enlightening environmental markets](#) (2020).
URL https://aemo.com.au/-/media/files/electricity/nem/planning_and_forecasting/inputs-assumptions-methodologies/2020/green-energy-markets-der-forecast-report.pdf?la=en
- [4] F. Brady, [A dictionary on electricity a joint project of cigre \(the international conference on large high voltage electrical systems\) and ahef \(the association for the history of electricity in france\)](#) (1996).
URL https://www.ewh.ieee.org/r10/nsw/subpages/history/electricity_in_australia.pdf
- [5] S. Jobson, [Power for the people: A history of electricity in sydney - on line opinion - 25/8/2004](#) (08 2004).
URL <https://www.onlineopinion.com.au/view.asp?article=2485&page=0>
- [6] R. Bentley, [Oil forecasts, past and present](#), *Energy Exploration and Exploitation* 20 (2002) 481–491. doi:[10.1260/014459802321615108](https://doi.org/10.1260/014459802321615108).
- [7] S. Smith, [Electricity and privatisation](#) (09 1997).
URL <https://www.parliament.nsw.gov.au/researchpapers/Documents/electricity-and-privatisation/ElectandPriv.pdf>
- [8] G. Parkinson, [Rooftop solar slashes demand levels and emissions across main grid](#) (08 2019).
URL <https://reneweconomy.com.au/rooftop-solar-slashes-demand-levels-and-emissions-across-main-grid>
- [9] D. Ürge Vorsatz, L. F. Cabeza, S. Serrano, C. Barreneche, K. Petrichenko, [Heating and cooling energy trends and drivers in buildings](#), *Renewable and Sustainable Energy Reviews* 41 (2015) 85–98. doi:<https://doi.org/10.1016/j.rser.2014.08.039>.
URL <https://www.sciencedirect.com/science/article/pii/S1364032114007151>

- [10] P. Pawar, Predicting hourly energy consumption of san diego (short and long term forecasts)— ii (06 2020).
URL <https://towardsdatascience.com/part-2-time-series-analysis-predicting->
- [11] J. . Griffin, Methodological advances in energy modelling: 1970-1990 (1993).
URL <https://www-jstor-org.wwwproxy1.library.unsw.edu.au/stable/pdf/41322485.pdf?refreqid=excelsior%3A93b2adff6bd9ecc4758484b4624cce7e>
- [12] G. Marcjasz, J. Lago, W. Rafal, Neural networks in day-ahead electricity price forecasting: single vs multiple outputs.
URL <https://arxiv.org/pdf/2008.08006.pdf>
- [13] R. C. Team, R: A language and environment for statistical computing, 2017.
URL <https://www.R-project.org/>
- [14] Y. Xie, J. Allaire, G. Grolemond, R markdown, the definitive guide, Chapman and Hall/CRC, 2018.
URL <https://bookdown.org/yihui/rmarkdown/>
- [15] AEMO, Electricity demand forecasting methodology information paper.
URL https://www.aemo.com.au/-/media/Files/Electricity/NEM/Planning_and_Forecasting/NEM_ES00/2019/Electricity-Demand-Forecasting-Methodology-Information-Paper.pdf
- [16] D. T. A. , Australian public holidays dates machine readable dataset, data.gov.au.
URL <https://data.gov.au/data/dataset/australian-holidays-machine-readable->
- [17] A. P. I. A. S. M. , Australian photovoltaic institute • pv performance by climate region (05 2021).
URL <https://pv-map.apvi.org.au/performance#4/>
- [18] T. Brinsmead, J. Hayward, P. Graham, Energy flagship australian electricity market analysis report to 2020 and 2030 final draft (2014).
URL <https://arena.gov.au/assets/2017/02/CSIRO-Electricity-market-analysis-for-IGEG.pdf>
- [19] <https://www.facebook.com/MachineLearningMastery>, Why one-hot encode data in machine learning? (07 2017).
URL <https://machinelearningmastery.com/why-one-hot-encode-data-in-machine-learning/>
- [20] ABCB, Climate zone map: New south wales and australian capital territory — australian building codes board (2015).
URL <https://www.abcb.gov.au/Resources/Tools-Calculators/Climate-Zone-Map-NSW-and-ACT>
- [21] K. P. Vatcheva, M. Lee, Multicollinearity in regression analyses conducted in epidemiologic studies, Epidemiology: Open Access 06. doi:10.4172/2161-1165.1000227.
URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4888898/>
- [22] A review of deep learning methods applied on load forecasting. doi:10.1109/ICMLA.2017.0-110.
- [23] c. o. o. A. o. B. o. Statistics, Australian bureau of statistics web site.
URL <https://www.abs.gov.au/websitedbs/d3310114.nsf/home/time+series+analysis:+the+basics#:~:text=WHAT%20IS%20A%20TIME%20SERIES>

- [24] R. Hyndman, G. Athanasopoulos, [8.7 ARIMA modelling in R — Forecasting: Principles and Practice \(2nd ed\)](#), Monash University, Australia, 2018.
URL <https://otexts.com/fpp2/arima-r.html#arima-r>
- [25] S. Ozden, M. Dursun, A. Aksöz, A. Saygın, Prediction and modelling of energy consumption on temperature control for greenhouses, Journal of Polytechnic [doi:10.2339/politeknik.417757](#).
- [26] J. Mesa Jiménez, L. Stokes, C. Moss, Q. Yang, V. N. Livina, Modelling energy demand response using long short-term memory neural networks, Energy Efficiency 13 (2020) 1263–1280. [doi:10.1007/s12053-020-09879-z](#).
- [27] S. A. Mousavi Maleki, H. Hizam, C. Gomes, [Estimation of hourly, daily and monthly global solar radiation on inclined surfaces: Models re-visited](#), Energies 10 (1). [doi:10.3390/en10010134](#).
URL <https://www.mdpi.com/1996-1073/10/1/134>

Appendix

Glossary

Item	Description
AEMO	Australian Energy Market Operator
VIF	Variance inflation factor
RMSE	Root Mean Square Error (We avoid RMSEP to avoid confusion)
PV	Photovoltaics
MLPRegressor	Multilayer Perceptron Regressor
LSTM	Long Short-Term Memory
XGBoost	XGBoost is a decision-tree-based ensemble Machine Learning algorithm
ARIMA	Autoregressive Integrated Moving Average
EDA	Exploratory Data Analysis
PowerBI	Power BI is a business analytics service by Microsoft
NSW	New South Wales
NEM	National Electricity Market
APVI	Australian PV Institute
kWp	kilowatts peak - peak power of a PV system or panel
NEMS	National Electricity Market System
HourMinute	A variable for the combination of Hour and Minute of 24 Hours time
DayName	Name of the seven week days
Spot market	Where financial instruments, such as commodities, currencies, and securities, are traded for immediate delivery
Box-Jenkins method	A Model is a mathematical model designed to forecast data ranges based on inputs from a specified time series
AIC	The Akaike information criterion
Adam	Adaptive Moment Estimation
SGD	Stochastic gradient descent
LBFGS	Limited-memory Broyden-Fletcher-Goldfarb-Shanno Algorithm
SARIMA	Seasonal Autoregressive Integrated Moving Average
ARIMAX	Autoregressive-moving-average model with exogenous inputs
BOM	Bureau of Meteorology
UV	Ultraviolet

Tables

VIF Factor	features
92.978786	FORECASTDEMAND
104.845037	YEAR
7.94458	MONTH
4.179661	DAY
6.746915	HOURL
2.75231	MINUTE
1.093466	PUBLICHOLIDAYS
1.572454	PVGeneration(MW)
1.074539	HOURLMINUTE_14-30
1.080066	HOURLMINUTE_16-30
1.12683	HOURLMINUTE_18-0
1.118781	HOURLMINUTE_18-30
1.134109	HOURLMINUTE_19-0
1.107565	HOURLMINUTE_2-0
1.118149	HOURLMINUTE_2-30
1.132064	HOURLMINUTE_20-0
1.116615	HOURLMINUTE_20-30
1.138781	HOURLMINUTE_21-0
1.081735	HOURLMINUTE_9-0
1.100514	HOURLMINUTE_9-30
1.999137	DAY_NAME_Friday
2.02014	DAY_NAME_Monday
2.079947	DAY_NAME_Saturday
2.138461	DAY_NAME_Sunday
1.997267	DAY_NAME_Thursday
2.001941	DAY_NAME_Tuesday
3.391713	SEASON_autumn
3.18336	SEASON_summer
2.920135	SEASON_winter
1.052842	TEMPERATURERANGES_extremelyhigh
1.117071	TEMPERATURERANGES_veryhigh
1.486606	TEMPERATURERANGES_verylow

Additional graphs




Figure 7.1: Expansion of the rooftop solar from 2012 to 2020




Figure 7.2: PV Performance comparison between Canterbury and Sutherland generated by single 5Kwp system




Figure 7.3: Outlier prediction by XGBoost highlighting the error for the event on the 4th of January